

Homework 1

by: Alexandre Adam
Collaborators: Olivier Vincent,
Ronan Legin

September 22, 2020

1 Probability and independence

a) Decomposition

We aim to validate

$$(X \perp Y, W \mid Z) \implies (X \perp Y \mid Z) \quad (1.1)$$

Proof. We suppose the statement $(X \perp Y, W \mid Z)$ is true. It follows from the definition of the conditional independence that $p(x, y, w \mid z) = p(x \mid z)p(y, w \mid z)$ for all $x \in \Omega_x$, $(y, w) \in \Omega_y \times \Omega_w$ and $z \in \Omega_z$. We then consider the marginalize $p(x, y, w \mid z)$:

$$\begin{aligned} p(x, y \mid z) &= \sum_{w \in \Omega_w} p(x, y, w \mid z) \\ &= \sum_{w \in \Omega_w} p(x \mid z)p(y, w \mid z) \\ &= p(x \mid z) \sum_{w \in \Omega_w} p(y, w \mid z) \\ &= p(x \mid z)p(y \mid z) \end{aligned}$$

from which we conclude that $(X \perp Y \mid Z) \quad \square$. By symmetry of the argument, we can show that $(X \perp W \mid Z)$ is true as well.

b)

We aim to validate

$$(X \perp Y \mid Z) \text{ and } (X, Y \perp W \mid Z) \implies (X \perp W \mid Z) \quad (1.2)$$

Proof. Suppose $(X, Y \perp W \mid Z)$ and $(X \perp Y \mid Z)$ are true. We know from the symmetry and decomposition properties of the conditional independence that $(X, Y \perp W \mid Z) \implies (W \perp X, Y \mid Z) \implies (X \perp W \mid Z)$. Therefore $(X \perp W \mid Z)$ is true \square .

c)

We aim to validate

$$(X \perp Y, W \mid Z) \text{ and } (Y \perp W \mid Z) \implies (X, W \perp Y \mid Z) \quad (1.3)$$

Proof. Suppose $(X \perp Y, W \mid Z)$ is true. Then it follows from the definition of conditional independence that

$$p(x, y, w \mid z) = p(x \mid z)p(y, w \mid z)$$

Then assume $(Y \perp W \mid Z)$ is true. The second factor can be factorized

$$p(x, y, w \mid z) = p(x \mid z)p(y \mid z)p(w \mid z)$$

From the decomposition property, we know $(X \perp W \mid Z)$ is true. Thus

$$p(x, y, w \mid z) = p(x, w \mid z)p(y \mid z)$$

From which we conclude $(X, W \perp Y \mid Z)$ is true \square .

d)

We aim to validate

$$(X \perp Y \mid Z) \text{ and } (X \perp Y \mid W) \implies (X \perp Y \mid Z, W) \quad (1.4)$$

Counter example. We consider the following situation: let Ω be a set of three identical elements $\Omega = \{1, 1, 1\}$. Let X, Y, W, Z be the action of removing an element or not, each independently distributed, but not necessarily identically distributed. Since the elements of Ω are identical, then any three R.V. are mutually and conditionally independent. It is not possible, however, to say the same about four R.V. since it is possible for 3 R.V. to change the distribution of the forth if the 3 given R.V. all choose to pick an element from the set. The probability distribution of the forth R.V. collapses into a delta function since only one choice remains. Therefore,

$$(X \perp Y \mid Z) \text{ and } (X \perp Y \mid Z) \not\implies (X \perp Y \mid Z, W)$$

2 Bayesian inference and MAP

Let $\mathbf{X}_1, \dots, \mathbf{X}_n \mid \boldsymbol{\pi} \stackrel{\text{iid}}{\sim} \text{Multinomial}(1, \boldsymbol{\pi})$ on k element. The values are sampled from a set of cardinality 2, that is $x_j^{(i)} \in \{0, 1\}$. Each R.V. has only one non-zero entry for a given trial, that is $\sum_{j=1}^k x_j^{(i)} = 1$.

We assume a Dirichlet prior $\boldsymbol{\pi} \sim \text{Dir}(\boldsymbol{\alpha})$ with a PDF

$$p(\boldsymbol{\pi} \mid \boldsymbol{\alpha}) = \frac{\Gamma(\sum_{i=1}^k \alpha_i)}{\prod_{j=1}^k \Gamma(\alpha_j)} \prod_{j=1}^k \pi_j^{\alpha_j - 1}$$

a)

Since the data is IID, they are mutually independent by definition. Being given the parameters of their Multinomial distribution (or a subset for that matter) does not change the independence of the \mathbf{X} 's. Thus,

$$(\mathbf{X}_i \perp \mathbf{X}_j \mid \boldsymbol{\pi}) \forall (i, j) \in \{1, \dots, k\} \times \{1, \dots, k\}$$

Of course, none of the vector can be mutually nor conditionally independent to $\boldsymbol{\pi}$ since it contains information about the distribution of the one hot vectors \mathbf{X}_i . In this case $\boldsymbol{\pi}$ are the probabilities of one of the k entry to be equal to one. Even giving one of these away is enough to impact the posterior distribution of the conditional $p(x_i \mid x_\ell, \pi_j)$ for example.

b)

The posterior distribution $p(\boldsymbol{\pi} \mid x_1, \dots, x_n)$ is computed via the Bayes rule

$$p(\boldsymbol{\pi} \mid \mathbf{x}_{1:n}) = \frac{p(\mathbf{x}_{1:n} \mid \boldsymbol{\pi})p(\boldsymbol{\pi})}{p(\mathbf{x}_{1:n})}$$

where $p(\boldsymbol{\pi}) = p(\boldsymbol{\pi} \mid \boldsymbol{\alpha})$ is the prior for $\boldsymbol{\pi}$ defined above. For the sake of determining the posterior distribution, we can postpone the derivation of the marginal likelihood. Therefore, the posterior must be

$$p(\boldsymbol{\pi} \mid \mathbf{x}_{1:n}) \propto \prod_{i=1}^n \prod_{j=1}^k \pi_j^{x_j^{(i)}} \prod_{\ell=1}^k \pi_\ell^{\alpha_\ell - 1}$$

We use the fact that we can swap around product operator for real numbers.

$$p(\boldsymbol{\pi} \mid \mathbf{x}_{1:n}) \propto \prod_{i=1}^n \prod_{j=1}^k \prod_{\ell=1}^k \pi_\ell^{\alpha_\ell - 1} \pi_j^{x_j^{(i)}} = \prod_{j=1}^k \pi_j^{\sum_{i=1}^n x_j^{(i)} + \alpha_j - 1}$$

We can readily see that the resulting distribution will be a Dirichlet with updated α_ℓ 's.

The posterior distribution is a Dirichlet distribution with parameters $\alpha'_j = \alpha_j + \sum_{i=1}^n x_j^{(i)}$.

c) Marginal Likelihood

The marginal likelihood $p(\mathbf{x}_{1:n})$ is a normalizing constant defined as the integral of the numerator (in the Bayes rule) over all instantiation of $\boldsymbol{\pi}$

$$p(\mathbf{x}_{1:n}) = \int_{\Delta_k} p(\mathbf{x}_{1:n} | \boldsymbol{\pi}) p(\boldsymbol{\pi}) d^{(k)} \boldsymbol{\pi}$$

where Δ_k is the probability simplex. In term of the quantities defined above, this is

$$p(\mathbf{x}_{1:n}) = \int_{\Delta_k} d^{(k)} \boldsymbol{\pi} \prod_{j=1}^k \pi_j^{\sum_{i=1}^n x_j^{(i)}} \left(\frac{\Gamma(\sum_{\ell=1}^k \alpha_\ell)}{\prod_{\ell=1}^k \Gamma(\alpha_\ell)} \prod_{\ell=1}^k \pi_\ell^{\alpha_\ell - 1} \right)$$

The π_j 's are independent variables since the simplex Δ_k is crucially defined as an affine plane in an Euclidean space which is supported by a set of orthonormal vectors. To evaluate this, we use the fact that the marginalized conjugate prior must sum to 1

$$\frac{\Gamma(\sum_{\ell=1}^k \alpha_\ell)}{\prod_{\ell=1}^k \Gamma(\alpha_\ell)} \int_{\Delta_k} d^{(k)} \boldsymbol{\pi} \prod_{j=1}^k \pi_j^{\alpha_j - 1} = 1$$

Thus, since both integral have the same form we assume

$$\int_{\Delta_k} d^{(k)} \boldsymbol{\pi} \pi_j^{\sum_{i=1}^n x_j^{(i)} + \alpha_j - 1} = \frac{\prod_{j=1}^k \Gamma(\sum_{i=1}^n x_j^{(i)} + \alpha_j)}{\Gamma(\sum_{j=1}^k (\sum_{i=1}^n x_j^{(i)} + \alpha_j))}$$

We then get

$$p(\mathbf{x}_{1:n}) = \frac{\Gamma(\sum_{\ell=1}^k \alpha_\ell)}{\prod_{\ell=1}^k \Gamma(\alpha_\ell)} \frac{\prod_{j=1}^k \Gamma(\sum_{i=1}^n x_j^{(i)} + \alpha_j)}{\Gamma(\sum_{j=1}^k (\sum_{i=1}^n x_j^{(i)} + \alpha_j))}$$

We notice that the first factor will cancel the one coming from the numerator in the posterior, and the second factor is the updated normalization factor of the Dirichlet.

d) $\hat{\boldsymbol{\pi}}_{\text{MAP}}$

The maximum *a posteriori* of the Multinomial distribution can be written in term of the log posterior

$$\hat{\boldsymbol{\pi}}_{\text{MAP}} \equiv \underset{\boldsymbol{\pi} \in \Delta_k}{\operatorname{argmax}} \log p(\boldsymbol{\pi} | \mathbf{x}_{1:n})$$

Where the probability simplex is defined as

$$\Delta_k = \left\{ \boldsymbol{\pi} \in \mathbb{R}^k \mid \pi_j \in [0, 1] \text{ and } \sum_{j=1}^k \pi_j = 1 \right\}$$

We define the constraint as $g(\boldsymbol{\pi}) = 1 - \sum_{j=1}^k \pi_j$. We notice that

$$\log p(\boldsymbol{\pi} | \mathbf{x}_{1:n}) = C + \sum_{j=1}^k \left(\sum_{i=1}^n x_j^{(i)} + \alpha_j - 1 \right) \log \pi_j$$

where C is the normalization constant. The optimisation of the log posterior becomes

$$\hat{\boldsymbol{\pi}}_{\text{MAP}} = \underset{(\boldsymbol{\pi}, \lambda) \in \mathbb{R}^{k+1}}{\operatorname{argmax}} \sum_{j=1}^k \left(\sum_{i=1}^n x_j^{(i)} + \alpha_j - 1 \right) \log \pi_j + \lambda g(\boldsymbol{\pi})$$

Here we ignore the normalizing constants which become an additive constants in the log posterior optimization problem. The solution is found where

$$\begin{aligned} \nabla_{\boldsymbol{\pi}} \log p(\boldsymbol{\pi} | \mathbf{x}_{1:n}) + \lambda g(\boldsymbol{\pi}) &= 0 \\ g(\boldsymbol{\pi}) &= 0 \end{aligned}$$

The first condition yields

$$\left[\nabla_{\boldsymbol{\pi}} \log p(\boldsymbol{\pi} | \mathbf{x}_{1:n}) + \lambda g(\boldsymbol{\pi}) \right]_\ell \bigg|_{\substack{\pi_\ell = \pi_\ell^* \\ \lambda = \lambda^*}} = 0 \implies \frac{\sum_{i=1}^n x_\ell^{(i)} + \alpha_\ell - 1}{\pi_\ell^*} = \lambda^*$$

Replacing this result in the second condition, we get

$$1 - \sum_{j=1}^k \frac{\sum_{i=1}^n x_j^{(i)} + \alpha_j - 1}{\lambda^*} = 0 \implies \lambda^* = n + \sum_{j=1}^k \alpha_j - k$$

Where we swapped the sum over the $x_j^{(i)}$ and used the fact that \mathbf{x}_j are one hot vectors. Thus

$$(\hat{\pi}_{\text{MAP}})_j = \pi_j^* = \frac{\sum_{i=1}^n x_j^{(i)} + \alpha_j - 1}{n + \sum_{j=1}^k \alpha_j - k} \in [0, 1]$$

The maximum likelihood estimator is, on the other hand,

$$(\hat{\pi}_{\text{MLE}})_j = \frac{\sum_{j=1}^n x_j^{(i)}}{n}$$

In the regime of extremely large k , knowing that $\alpha_j > 1 \forall j$, we expect the sum $\sum_{j=1}^k \alpha_j - k \gg 1$ to become non-negligible. In turns, this means that we expect

$$(\hat{\pi}_{\text{MAP}})_j < (\hat{\pi}_{\text{MLE}})_j$$

3 Properties of estimators

a) Poisson

Let n trials $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} \text{Poisson}(\lambda)$ where $\lambda = \mathbb{E}_x[x]$. The pmf of the Poisson is

$$p(x | \lambda) = \frac{\lambda^x}{x!} e^{-\lambda} \forall x \in \mathbb{N}$$

Such that the pmf of n trials should be

$$p(x_{1:n} | \lambda) \propto \prod_{j=1}^n p(x_j | \lambda)$$

I MLE

Using the log likelihood, we define the MLE estimation of λ as

$$\hat{\lambda}_{\text{MLE}} = \underset{\lambda \in \mathbb{R}_{>0}}{\text{argmax}} \sum_{j=1}^n (x_j \log \lambda - \lambda)$$

Which is found where

$$\left. \nabla_{\lambda} \log p(x_{1:n} | \lambda) \right|_{\lambda=\lambda^*} = 0$$

That is

$$\nabla_{\lambda} \log p(x_{1:n} | \lambda) = \lambda^{-1} \sum_{j=1}^n (x_j - 1)$$

Thus

$$\hat{\lambda}_{\text{MLE}} = \frac{1}{n} \sum_{j=1}^n x_j$$

II Bias

The bias is defined as

$$\text{Bias}(\lambda, \hat{\lambda}_{\text{MLE}}) \equiv \mathbb{E}_x[\hat{\lambda}_{\text{MLE}}] - \lambda$$

The expectation value of the MLE estimator is

$$\begin{aligned}\mathbb{E}_x[\hat{\lambda}_{\text{MLE}}] &= \mathbb{E}_x \left[\frac{1}{n} \sum_{j=1}^n x'_j \right] \\ &= \frac{1}{n} \sum_{j=1}^k \mathbb{E}_x[x_j] \\ &= \lambda\end{aligned}$$

Therefore the **MLE estimator of a Poisson distribution is an unbiased estimator.**

III Variance

The variance of the estimator is

$$\text{Var}(\hat{\lambda}_{\text{MLE}}) \equiv \mathbb{E}_X[\hat{\lambda}_{\text{MLE}}^2] - \mathbb{E}_x^2[\hat{\lambda}_{\text{MLE}}]$$

We need to evaluate the first term. To do this, we first use the Multinomial theorem to expand $\hat{\lambda}_{\text{MLE}}^2$:

$$\hat{\lambda}_{\text{MLE}}^2 = \frac{1}{n^2} \sum_{k_1 + \dots + k_n = 2} \binom{2}{k_1, \dots, k_n} x_1^{k_1} \dots x_n^{k_n}$$

Then we use both the linearity of the expectation operator and the fact that the R.V. X_1, \dots, X_n are independent to factorize the expectation of a "cross" product

$$\mathbb{E}_x[X_i X_j] = \mathbb{E}_x[X_i] \mathbb{E}_x[X_j], \quad \forall i, j \in \{\text{iid}\}$$

to get

$$\mathbb{E}_x[\hat{\lambda}_{\text{MLE}}^2] = \frac{1}{n^2} \sum_{k_1 + \dots + k_n = 2} \binom{2}{k_1, \dots, k_n} \mathbb{E}_x[x_1^{k_1}] \dots \mathbb{E}_x[x_n^{k_n}]$$

The sum can be separated into quadratic and linear term, s.t.

$$\mathbb{E}_x[\hat{\lambda}_{\text{MLE}}^2] = \frac{1}{n^2} \left(n \mathbb{E}_x[x^2] + 2 \binom{n}{2} \lambda^2 \right)$$

We used the fact that $\mathbb{E}_x[1] = 1$ and $\mathbb{E}_x[x_j] = \lambda$, $\forall j$. To estimate the quadratic term, we can use a magic trick by adding zero inside the operator argument. Using its linear property

$$\mathbb{E}_x[x^2] = \mathbb{E}_x[x(x-1) + x] = \mathbb{E}_x[x(x-1)] + \lambda$$

It turns out that

$$\begin{aligned}\mathbb{E}_x[x(x-1)] &= \sum_{x=0}^{\infty} x(x-1) \frac{\lambda^x}{x!} e^{-\lambda} \\ &= e^{-\lambda} \sum_{x=2}^{\infty} \frac{\lambda^x}{(x-2)!} \\ &= e^{-\lambda} \sum_{x=0}^{\infty} \frac{\lambda^{x+2}}{x!} \\ &= \lambda^2\end{aligned}$$

By noticing the sum is the Taylor series of e^λ . In the end, we get

$$\mathbb{E}_x[\hat{\lambda}_{\text{MLE}}^2] = \frac{1}{n^2} (n\lambda + n^2\lambda^2)$$

Where we expanded the binom coefficient $2 \binom{n}{2} = n(n-1)$. The variance is thus

$$\boxed{\text{Var}(\hat{\lambda}_{\text{MLE}}) = \frac{\lambda}{n}}$$

IV Consistency

As $n \rightarrow \infty$, the estimator give an unbiased estimate of λ with a variance that goes to 0. Thus, the **estimator is consistent.**

b) Bernoulli

Let $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} \text{Bernoulli}(p)$ and let $n > 10$. We consider the estimator

$$\hat{p} \equiv \frac{1}{10} \sum_{i=1}^{10} X_i$$

I Bias

We first note that the expected value of a Bernoulli is

$$\mathbb{E}_x[x] = p$$

Since $x \in \{0, 1\}$ and p is the probability that $X = 1$. Therefore,

$$\text{Bias}(p, \hat{p}) = \frac{1}{10} \sum_{j=1}^{10} \mathbb{E}_x[x_j] - p = 0$$

\hat{p} is an unbiased estimator.

II Variance

The variance is

$$\begin{aligned} \text{Var}(\hat{p}) &= \mathbb{E}_x[\hat{p}^2] - \mathbb{E}_x[\hat{p}]^2 \\ &= \frac{1}{100} (10 \mathbb{E}_x[x^2] + 90 \mathbb{E}_x[x]) - p^2 \\ &= \frac{p}{10} + p^2 \left(\frac{90}{100} - 1 \right) \\ &= \frac{1}{10} (p - p^2) \end{aligned}$$

III Consistency

This estimator is not consistent since the variance is constant as $n \rightarrow \infty$.

c) Uniform

Let $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} \text{Uniform}(0, \theta)$. The pdf of this distribution is

$$p(x_i | \theta) = \begin{cases} \theta^{-1}, & x \in [0, \theta] \\ 0, & \text{otherwise} \end{cases} = \frac{1}{\theta} \mathbf{1}_{\{0 \leq x_i \leq \theta\}}$$

for $\theta \in \mathbb{R}_{>0}$. We used the indicator function

$$\mathbf{1}_A(x) \equiv \begin{cases} 1, & x \in A \\ 0, & x \notin A \end{cases}$$

I MLE

Given n samples, we want an estimator of the maximum possible value of \mathbf{X} . The MLE is

$$\hat{\theta}_{\text{MLE}} = \underset{\theta \in \mathbb{R}_{>0}}{\text{argmax}} \frac{1}{\theta^n} \prod_{i=1}^n \mathbf{1}_{\{0 \leq x_i \leq \theta\}}$$

Where we used the fact that the data is iid. We can see that the product only depends on the boundary cases of the dataset, that is

$$\hat{\theta}_{\text{MLE}} = \underset{\theta \in \mathbb{R}_{>0}}{\text{argmax}} \frac{1}{\theta^n} \mathbf{1}_{\{0 \leq \min \mathbf{X}\}} \mathbf{1}_{\{\max \mathbf{X} \leq \theta\}}$$

One can see that θ should be as low as possible to maximize θ^{-n} , yet not too low s.t. it make the second indicator function 0. The obvious choice is therefore

$$\hat{\theta}_{\text{MLE}} = \max \mathbf{X}$$

We show that $T(\theta) = \hat{\theta}_{\text{MLE}}$ is a sufficient statistic. To show this, we use the Fisher-Neyman theorem which guarantees that the statistic is sufficient if the probability density can be factorized as $p(\mathbf{X}) = h(\mathbf{X})g(\theta, T(\theta))$. First, we use the fact that the data is iid:

$$p(\mathbf{X}) = \prod_{i=1}^n p(x_i)$$

Then replacing by the Uniform distribution

$$p(\mathbf{X}) = \prod_{i=1}^n \frac{1}{\theta} \mathbf{1}_{\{0 \leq x_i \leq \theta\}}$$

For the probability to be non-zero, only the boundary cases are important. That is

$$p(\mathbf{X}) = \frac{1}{\theta^n} \mathbf{1}_{\{0 \leq \min \mathbf{X}\}} \mathbf{1}_{\{\max \mathbf{X} \leq \theta\}} = h(\mathbf{X}) \frac{1}{\theta^n} \mathbf{1}_{\{T(\theta) \leq \theta\}}$$

We identify $h(\mathbf{X}) = \mathbf{1}_{\{0 \leq \min \mathbf{X}\}}$ and the rest with the function $g(\theta, T(\theta))$. Thus we have shown $T(\theta)$ is a sufficient statistic by the Fisher-Neyman theorem.

II Bias

The bias of this estimator is

$$\text{Bias}(\theta, \hat{\theta}_{\text{MLE}}) = \mathbb{E}_c[\hat{\theta}_{\text{MLE}}] - \theta$$

Where $c = \max \mathbf{X}$. To compute the expectation value of the MLE estimator, we must first compute the pdf with respect to c . We notice that the likelihood of the maximum in the set \mathbf{X} to be smaller than c is

$$p(\max \mathbf{X} < c \mid \theta) = \prod_{i=1}^n \left(\frac{c}{\theta}\right) = \left(\frac{c}{\theta}\right)^n$$

from the hint. To get the pdf, we derive this expression with respect to c (the likelihood over the region $x_i \in [0, c]$ is an integral, so we expect the pdf to be the derivative of this integral):

$$p(\max \mathbf{X} = c \mid \theta) = n \frac{c^{n-1}}{\theta^n}$$

Therefore,

$$\mathbb{E}_c[\hat{\theta}_{\text{MLE}}] = n \int_0^\theta c \frac{c^{n-1}}{\theta^n} dc = \frac{n}{n+1} \theta$$

And we get

$$\left\| \text{Bias}(\theta, \hat{\theta}_{\text{MLE}}) \right\|_2 = \frac{\theta}{n+1}$$

III Variance

The variance is

$$\text{Var}(\hat{\theta}_{\text{MLE}}) = \mathbb{E}_c[\hat{\theta}_{\text{MLE}}^2] - \mathbb{E}_c^2[\hat{\theta}_{\text{MLE}}]$$

Thus

$$\text{Var}(\hat{\theta}_{\text{MLE}}) = \theta^2 \left(\frac{n}{n+2} - \frac{n^2}{(n+1)^2} \right)$$

IV Consistency

The estimator is consistent because the bias and the variance go to zero as $n \rightarrow \infty$

d) Gaussian

Let $X_1, \dots, X_n \stackrel{\text{iid}}{\sim} \mathcal{N}(\mu, \sigma^2)$, $\sigma, \mu \in \mathbb{R}$. We define the mean as $\bar{X} = \frac{1}{n} \sum_{i=1}^n x_i$.

I MLE

The likelihood, using the fact that the data is iid, is

$$p(\mathbf{X} \mid \mu, \sigma) = \frac{1}{(\sigma\sqrt{2\pi})^n} \prod_{i=1}^n \exp\left(-\frac{1}{2} \frac{(x_i - \mu)^2}{\sigma^2}\right)$$

The MLE for μ and σ^2 can be derived from the log likelihood

$$\hat{\theta}_{\text{MLE}} = (\hat{\mu}_{\text{MLE}}, \hat{\sigma}_{\text{MLE}}^2) = \underset{(\mu, \sigma^2) \in \mathbb{R}^2}{\operatorname{argmax}} -\frac{n}{2} \log \sigma^2 - \frac{1}{2} \sum_{i=1}^n \frac{(x_i - \mu)^2}{\sigma^2}$$

The solution is found where

$$\nabla_{\theta} \log p(\mathbf{X} \mid \theta) = 0$$

That is

$$\partial_{\mu} \implies \sum_{i=1}^n (x_i - \hat{\mu}) = 0$$

From which we find

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i = \bar{X}$$

Also, we have

$$\partial_{\sigma^2} \implies -\frac{n}{2\hat{\sigma}^2} + \frac{1}{2} \sum_{i=1}^n \frac{(x_i - \bar{X})^2}{\hat{\sigma}^4} = 0$$

Thus

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{X})^2$$

II Bias

The bias for the normal variance MLE $\hat{\sigma}_{\text{MLE}}^2$ is

$$\text{Bias}(\sigma^2, \hat{\sigma}_{\text{MLE}}^2) = \mathbb{E}_x[\hat{\sigma}_{\text{MLE}}^2] - \sigma^2$$

Thus,

$$\text{Bias}(\sigma^2, \hat{\sigma}_{\text{MLE}}^2) = \frac{1}{n} \sum_{i=1}^n (\mathbb{E}_x[x_i^2] - 2 \mathbb{E}_x[x_i] \mathbb{E}_x[\bar{X}] + \mathbb{E}_x[\bar{X}^2]) - \sigma^2$$

With the definition of the variance, we can replace $\mathbb{E}_x[x^2] = \sigma^2 + \mu^2$. Thus,

$$\text{Bias}(\sigma^2, \hat{\sigma}_{\text{MLE}}^2) = \frac{1}{n} \sum_{i=1}^n (\sigma^2 + \mu^2 - 2\mu^2 + \mathbb{E}_x[\bar{X}^2]) - \sigma^2$$

Expanding the square of the mean

$$\bar{X}^2 = \frac{1}{n^2} \sum_{i=1}^n \left(x_i^2 + 2 \sum_{j>i} x_i x_j \right)$$

Thus

$$\mathbb{E}_x[\bar{X}^2] = \frac{1}{n} \sigma^2 + \frac{1}{n} \mu^2 + \frac{(n-1)}{n} \mu^2 = \frac{1}{n} \sigma^2 + \mu^2$$

Finally,

$$\|\text{Bias}(\sigma^2, \hat{\sigma}_{\text{MLE}}^2)\|_2 = \frac{\sigma^2}{n}$$

III Variance

We can use the chi-squared distribution for which we know the variance

$$\text{Var}(\chi_{n-1}^2) = 2(n-1)$$

Knowing that

$$\chi_{n-1}^2 \equiv \frac{1}{\sigma^2} \sum_{i=1}^n (x_i - \bar{X})^2 = \frac{n\hat{\sigma}_{\text{MLE}}^2}{\sigma^2}$$

Thus

$$\text{Var}(\hat{\sigma}_{\text{MLE}}^2) = \text{Var}\left(\frac{\chi_{n-1}^2 \sigma^2}{n}\right) = \frac{2\sigma^4(n-1)}{n^2}$$

IV Consistency

Both the variance and the bias of the estimator go to 0 as $n \rightarrow \infty$, so the **estimator of the variance is consistent**.